

Image_Classification

May 26, 2024

```
[1]: import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow.keras import datasets, models
import tensorflow.keras.layers as tfl
from tensorflow.keras.utils import to_categorical
from tensorflow.math import confusion_matrix
from tensorflow.keras.utils import to_categorical
```

```
[2]: (train_X, train_Y), (test_X, test_Y) = datasets.cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 [=====] - 2s 0us/step


```
[3]: print(train_X.shape)
print(test_X.shape)
print(train_Y.shape)
print(test_Y.shape)
```

```
(50000, 32, 32, 3)
(10000, 32, 32, 3)
(50000, 1)
(10000, 1)
```


```
[4]: classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "↵
↵"horse", "ship", "truck"]

def plot_image(X,Y,index):
    plt.figure(figsize=(15,2))
    plt.imshow(X[index])
    plt.xlabel(classes[Y[index][0]])
```

```
[5]: plot_image(train_X, train_Y, 2)
```



```
[6]: plot_image(train_X, train_Y, 26)
```



```
[7]: train_X = train_X/255
test_X = test_X/255

train_Y = to_categorical(train_Y)
test_Y = to_categorical(test_Y)
```

```
[8]: ann = models.Sequential([
    tf1.Flatten(input_shape=(32,32,3)),
    tf1.Dense(3000, activation='relu'),
    tf1.Dense(1000, activation='relu'),
    tf1.Dense(10,activation='sigmoid')
])
```

```
[9]: ann.compile(optimizer='adam',
                loss='categorical_crossentropy',
                metrics=['accuracy'])
```

```
[10]: ann.fit(train_X, train_Y, epochs=5)
```

```
Epoch 1/5
1563/1563 [=====] - 11s 6ms/step - loss: 1.8758 -
accuracy: 0.3295
Epoch 2/5
1563/1563 [=====] - 9s 6ms/step - loss: 1.6785 -
accuracy: 0.3971
Epoch 3/5
1563/1563 [=====] - 10s 7ms/step - loss: 1.5995 -
accuracy: 0.4254
Epoch 4/5
1563/1563 [=====] - 10s 7ms/step - loss: 1.5456 -
accuracy: 0.4455
Epoch 5/5
1563/1563 [=====] - 11s 7ms/step - loss: 1.5117 -
accuracy: 0.4599
```

```
[10]: <keras.src.callbacks.History at 0x78955bcaa1d0>
```

```
[11]: ann.evaluate(test_X,test_Y)
```

```
313/313 [=====] - 1s 4ms/step - loss: 1.5255 -
accuracy: 0.4623
```

```
[11]: [1.5255182981491089, 0.46230000257492065]
```

```
[12]: pred_test_Y = ann.predict(test_X)
```

313/313 [=====] - 1s 2ms/step

```
[13]: conf_mat = confusion_matrix(test_Y.argmax(axis=1), pred_test_Y.argmax(axis=1))
```

```
[14]: import seaborn as sns
plt.figure(figsize=(8,8))
sns.heatmap(conf_mat, annot=True, fmt='.1f', cmap='Blues')
```

[14]: <Axes: >



```
[15]: from sklearn.metrics import confusion_matrix, classification_report
```

```
[16]: print(classification_report(test_Y.argmax(axis=1), pred_test_Y.argmax(axis=1) ))
```

	precision	recall	f1-score	support
0	0.58	0.46	0.51	1000
1	0.60	0.62	0.61	1000
2	0.30	0.47	0.36	1000
3	0.34	0.28	0.31	1000
4	0.37	0.40	0.38	1000
5	0.47	0.26	0.33	1000
6	0.40	0.62	0.49	1000

7	0.68	0.35	0.46	1000
8	0.57	0.65	0.61	1000
9	0.55	0.51	0.53	1000
accuracy			0.46	10000
macro avg	0.49	0.46	0.46	10000
weighted avg	0.49	0.46	0.46	10000

```
[17]: cnn = models.Sequential([

    # cnn
    tf1.Conv2D(filters=32, kernel_size=(3,3), activation='relu',
    ↪input_shape=(32,32,3)),
    tf1.MaxPooling2D((2,2)),

    tf1.Conv2D(filters=64, kernel_size=(3,3), activation='relu'),
    tf1.MaxPooling2D((2,2)),

    tf1.Conv2D(filters=128, kernel_size=(3,3), activation='relu'),
    tf1.MaxPooling2D((2,2)),

    # Dense
    tf1.Flatten(),
    tf1.Dense(128, activation='relu'),
    tf1.Dense(10, activation='softmax')

])
```

```
[18]: cnn.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
```

```
[19]: cnn.fit(train_X, train_Y, epochs=5)
```

```
Epoch 1/5
1563/1563 [=====] - 14s 6ms/step - loss: 1.4877 -
accuracy: 0.4564
Epoch 2/5
1563/1563 [=====] - 11s 7ms/step - loss: 1.0982 -
accuracy: 0.6124
Epoch 3/5
1563/1563 [=====] - 12s 7ms/step - loss: 0.9406 -
accuracy: 0.6711
Epoch 4/5
1563/1563 [=====] - 13s 8ms/step - loss: 0.8412 -
accuracy: 0.7042
Epoch 5/5
```

```
1563/1563 [=====] - 11s 7ms/step - loss: 0.7643 -  
accuracy: 0.7322
```

```
[19]: <keras.src.callbacks.History at 0x78955c694760>
```

```
[20]: cnn.evaluate(test_X, test_Y)
```

```
313/313 [=====] - 1s 4ms/step - loss: 0.9857 -  
accuracy: 0.6742
```

```
[20]: [0.9857380986213684, 0.6741999983787537]
```

```
[21]: pred_test_Y = cnn.predict(test_X)
```

```
313/313 [=====] - 1s 2ms/step
```


```
[22]: conf_matrix = confusion_matrix(test_Y.argmax(axis=1), pred_test_Y.argmax(axis=1))
```

```
[23]: print(conf_matrix)
```

```
[[700  25  20   6  41   1   2  12 119  74]  
 [ 13 785   0   3   8   0   3   3  41 144]  
 [ 89  12 361 102 228  25  45  56  38  44]  
 [ 22  14  18 554 132  41  36  62  51  70]  
 [ 12   6   8  45 775   5  29  69  31  20]  
 [ 11  11  19 266  92 371  17 120  46  47]  
 [  9  11  17  80 101   5 678  17  34  48]  
 [ 14   4  10  40  92  14   3 753   8  62]  
 [ 37  14   1   6   7   0   1   3 880  51]  
 [ 12 43   1   6   5   2   2   5  39 885]]
```

```
[24]: plt.figure(figsize=(8,8))  
      sns.heatmap(conf_matrix, annot=True, fmt='.1f', cmap='Blues')
```

```
[24]: <Axes: >
```



```
[25]: print(classification_report(test_Y.argmax(axis=1), pred_test_Y.argmax(axis=1)))
```

	precision	recall	f1-score	support
0	0.76	0.70	0.73	1000
1	0.85	0.79	0.82	1000
2	0.79	0.36	0.50	1000
3	0.50	0.55	0.53	1000
4	0.52	0.78	0.62	1000
5	0.80	0.37	0.51	1000
6	0.83	0.68	0.75	1000
7	0.68	0.75	0.72	1000
8	0.68	0.88	0.77	1000
9	0.61	0.89	0.72	1000
accuracy			0.67	10000
macro avg	0.70	0.67	0.67	10000
weighted avg	0.70	0.67	0.67	10000